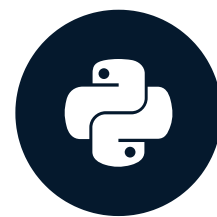


Two numeric explanatory variables

INTERMEDIATE REGRESSION WITH STATSMODELS IN PYTHON



Maarten Van den Broeck

Content Developer at DataCamp


Visualizing three numeric variables

- 3D scatter plot
- 2D scatter plot with response as color

Another column for the fish dataset


species	mass_g	length_cm	height_cm
Bream	1000	33.5	18.96
Bream	925	36.2	18.75
Roach	290	24.0	8.88
Roach	390	29.5	9.48
Perch	1100	39.0	12.80
Perch	1000	40.2	12.60
Pike	1250	52.0	10.69
Pike	1650	59.0	10.81

3D scatter plot



2D scatter plot, color for response

```
sns.scatterplot(x="length_cm",  
                y="height_cm",  
                data=fish,  
                hue="mass_g")
```



Modeling with two numeric explanatory variables

```
mdl_mass_vs_both = ols("mass_g ~ length_cm + height_cm",  
                        data=fish).fit()  
  
print(mdl_mass_vs_both.params)
```

```
Intercept    -622.150234  
length_cm     28.968405  
height_cm     26.334804
```

The prediction flow

```
from itertools import product

length_cm = np.arange(5, 61, 5)
height_cm = np.arange(2, 21, 2)

p = product(length_cm, height_cm)

explanatory_data = pd.DataFrame(p,
                                columns=["length_cm",
                                         "height_cm"])

prediction_data = explanatory_data.assign(
    mass_g = mdl_mass_vs_both.predict(explanatory_data))

print(prediction_data)
```

	length_cm	height_cm	mass_g
0	5	2	-424.638603
1	5	4	-371.968995
2	5	6	-319.299387
3	5	8	-266.629780
4	5	10	-213.960172
..
115	60	12	1431.971694
116	60	14	1484.641302
117	60	16	1537.310909
118	60	18	1589.980517
119	60	20	1642.650125

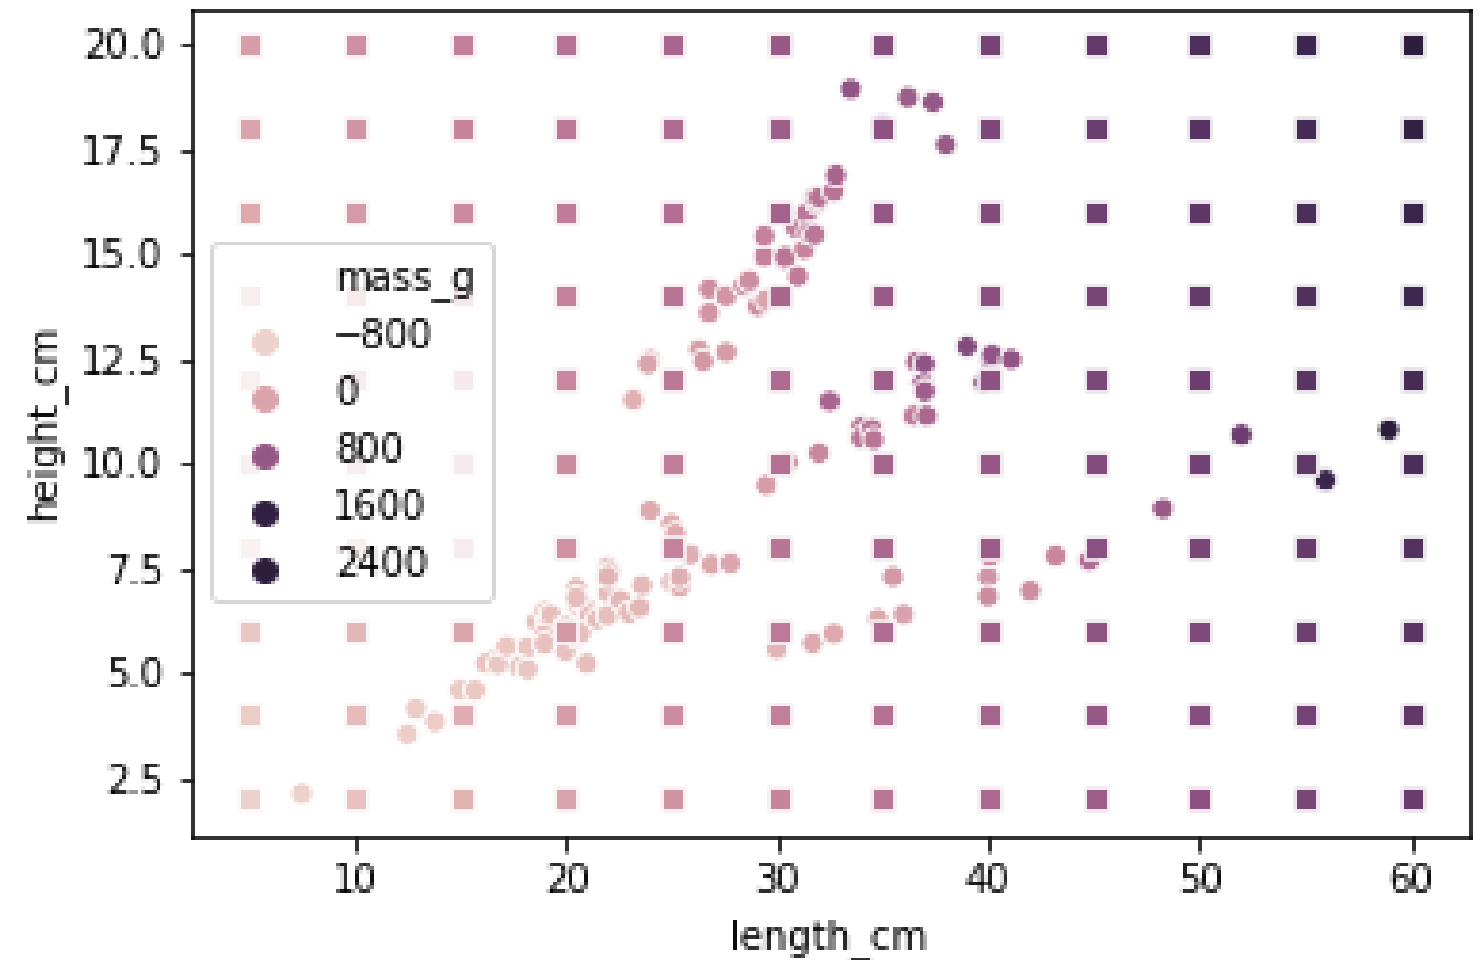
[120 rows x 3 columns]

Plotting the predictions

```
sns.scatterplot(x="length_cm",  
               y="height_cm",  
               data=fish,  
               hue="mass_g")
```

```
sns.scatterplot(x="length_cm",  
               y="height_cm",  
               data=prediction_data,  
               hue="mass_g",  
               legend=False,  
               marker="s")
```

```
plt.show()
```



Including an interaction

```
mdl_mass_vs_both_inter = ols("mass_g ~ length_cm * height_cm",  
                             data=fish).fit()
```

```
print(mdl_mass_vs_both_inter.params)
```

Intercept	159.107480
length_cm	0.301426
height_cm	-78.125178
length_cm:height_cm	3.545435

The prediction flow with an interaction

```
length_cm = np.arange(5, 61, 5)
height_cm = np.arange(2, 21, 2)

p = product(length_cm, height_cm)

explanatory_data = pd.DataFrame(p,
                                columns=["length_cm",
                                         "height_cm"])


prediction_data = explanatory_data.assign(
    mass_g = mdl_mass_vs_both_inter.predict(explanatory_data))
```

Plotting the predictions

```
sns.scatterplot(x="length_cm",  
               y="height_cm",  
               data=fish,  
               hue="mass_g")
```

```
sns.scatterplot(x="length_cm",  
               y="height_cm",  
               data=prediction_data,  
               hue="mass_g",  
               legend=False,  
               marker="s")
```

```
plt.show()
```




Let's practice!

INTERMEDIATE REGRESSION WITH STATSMODELS IN PYTHON

More than two explanatory variables

INTERMEDIATE REGRESSION WITH STATSMODELS IN PYTHON




Maarten Van den Broeck

Content Developer at DataCamp

From last time

```
sns.scatterplot(x="length_cm",  
                y="height_cm",  
                data=fish,  
                hue="mass_g")
```

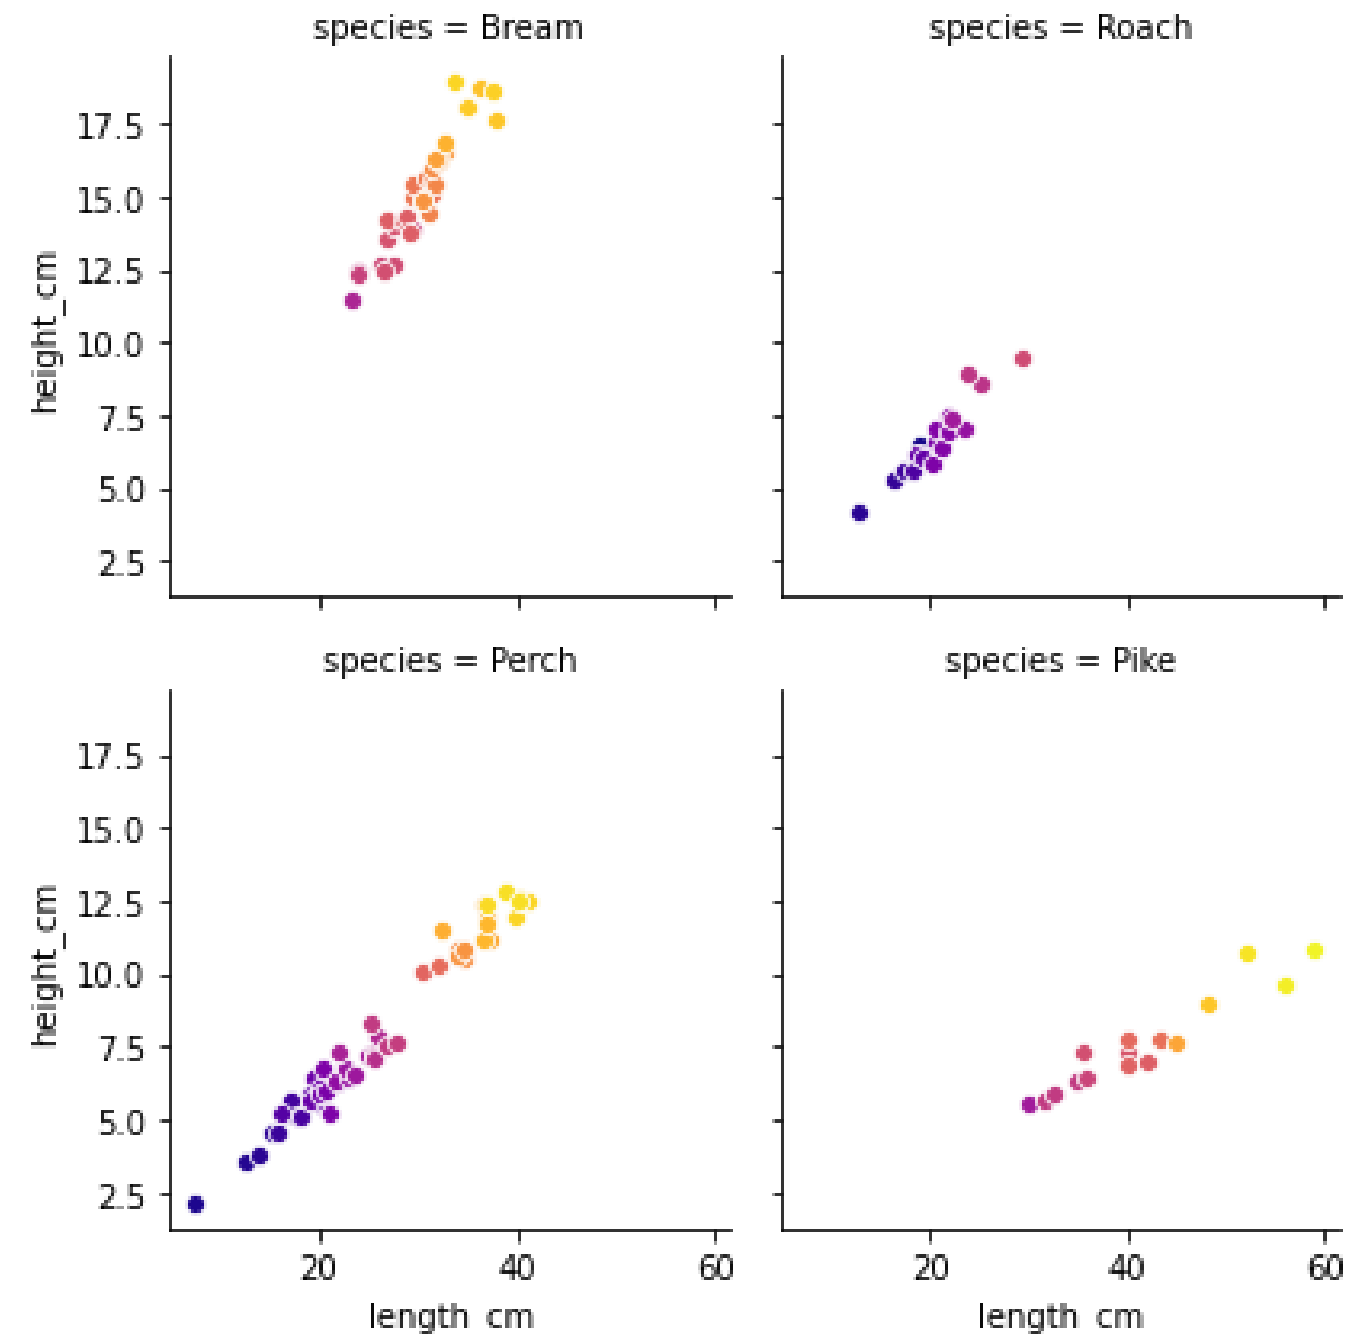


Faceting by species

```
grid = sns.FacetGrid(data=fish,  
                      col="species",  
                      hue="mass_g",  
                      col_wrap=2,  
                      palette="plasma")
```

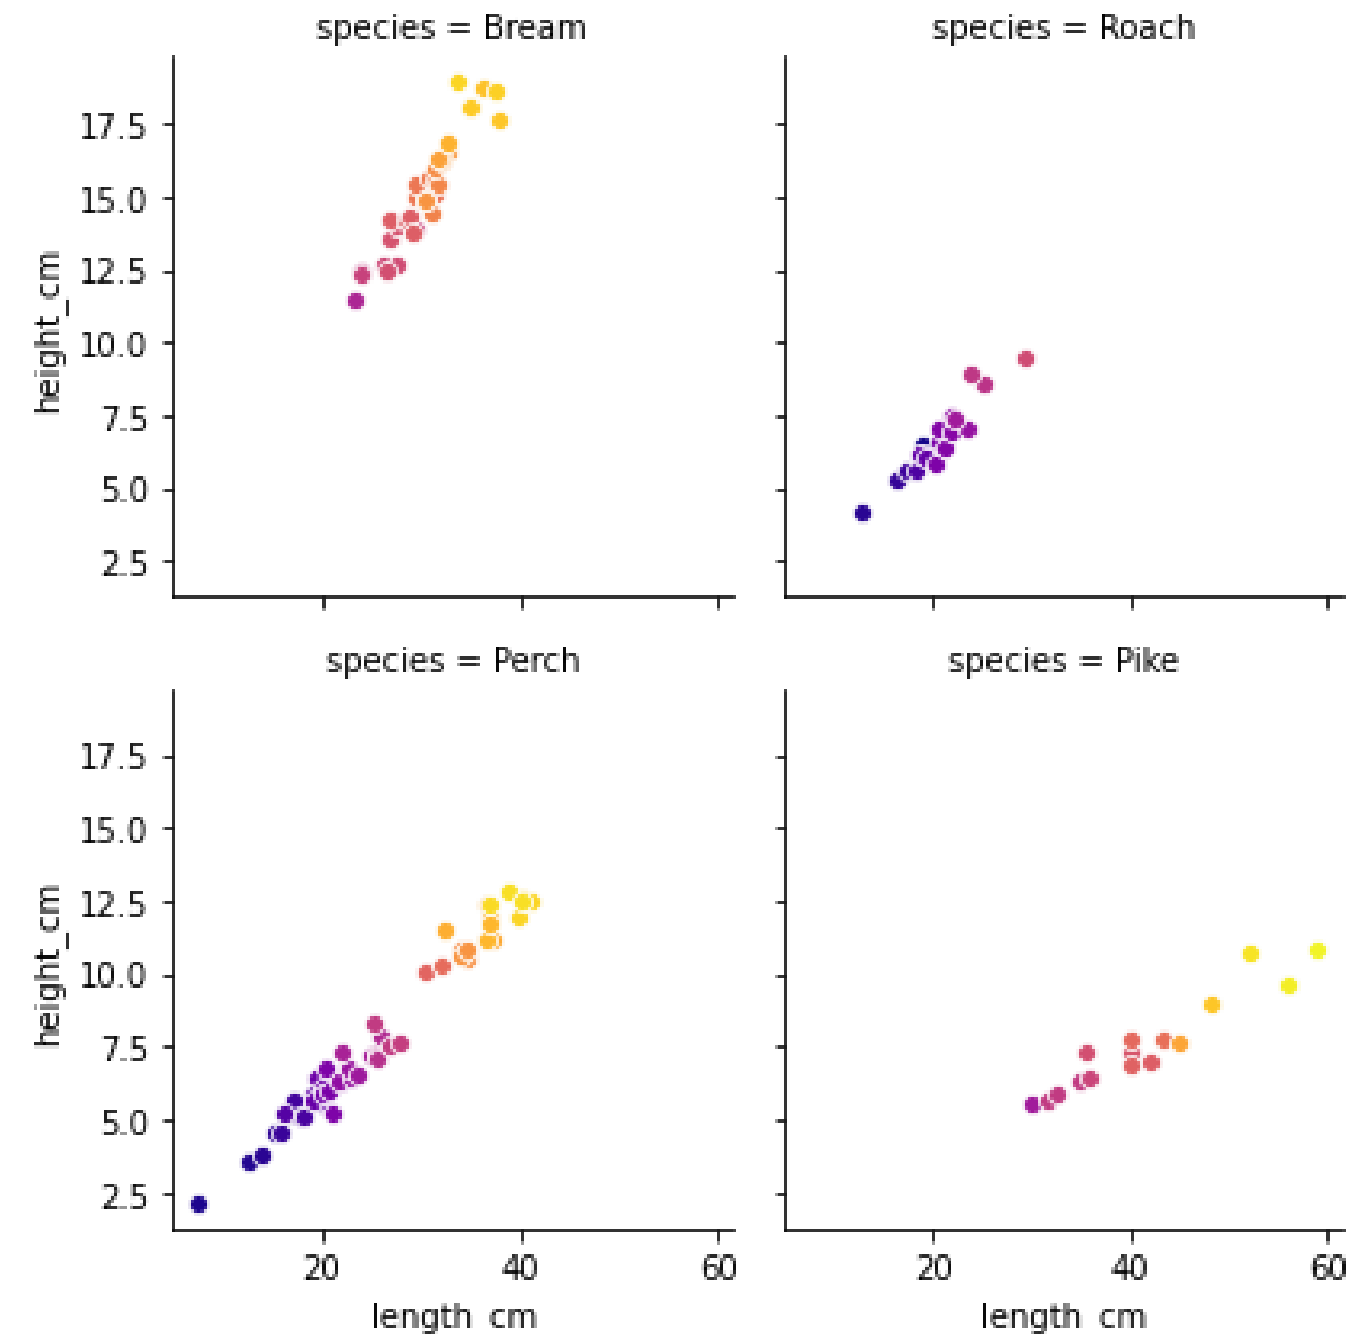
```
grid.map(sns.scatterplot,  
         "length_cm",  
         "height_cm")
```

```
plt.show()
```



Faceting by species

- It's possible to use more than one categorical variable for faceting
- Beware of faceting overuse
- Plotting becomes harder with increasing number of variables



Different levels of interaction

No interactions

```
ols("mass_g ~ length_cm + height_cm + species + 0", data=fish).fit()
```

two-way interactions between pairs of variables

```
ols(  
    "mass_g ~ length_cm + height_cm + species +  
    length_cm:height_cm + length_cm:species + height_cm:species + 0", data=fish).fit()
```

three-way interaction between all three variables

```
ols(  
    "mass_g ~ length_cm + height_cm + species +  
    length_cm:height_cm + length_cm:species + height_cm:species + length_cm:height_cm:species + 0", data=fish).fit()
```

All the interactions

```
ols(  
    "mass_g ~ length_cm + height_cm + species +  
    length_cm:height_cm + length_cm:species + height_cm:species + length_cm:height_cm:species + 0",  
    data=fish).fit()
```

same as

```
ols(  
    "mass_g ~ length_cm * height_cm * species + 0",  
    data=fish).fit()
```

Only two-way interactions

```
ols(  
    "mass_g ~ length_cm + height_cm + species +  
    length_cm:height_cm + length_cm:species + height_cm:species + 0",  
    data=fish).fit()
```

same as

```
ols(  
    "mass_g ~ (length_cm + height_cm + species) ** 2 + 0",  
    data=fish).fit()
```

The prediction flow

```
mdl_mass_vs_all = ols(
    "mass_g ~ length_cm * height_cm * species + 0",
    data=fish).fit()

length_cm = np.arange(5, 61, 5)
height_cm = np.arange(2, 21, 2)
species = fish["species"].unique()

p = product(length_cm, height_cm, species)

explanatory_data = pd.DataFrame(p,
                                columns=["length_cm",
                                         "height_cm",
                                         "species"])

prediction_data = explanatory_data.assign(
    mass_g = mdl_mass_vs_all.predict(explanatory_data))
```

	length_cm	height_cm	species	mass_g
0	5	2	Bream	-570.656437
1	5	2	Roach	31.449145
2	5	2	Perch	43.789984
3	5	2	Pike	271.270093
4	5	4	Bream	-451.127405
..
475	60	18	Pike	2690.346384
476	60	20	Bream	1531.618475
477	60	20	Roach	2621.797668
478	60	20	Perch	3041.931709
479	60	20	Pike	2926.352397

[480 rows x 4 columns]

Let's practice!

INTERMEDIATE REGRESSION WITH STATSMODELS IN PYTHON

How linear regression works


INTERMEDIATE REGRESSION WITH STATSMODELS IN PYTHON




Maarten Van den Broeck

Content Developer at DataCamp

The standard simple linear regression plot



Visualizing residuals



A metric for the best fit

The simplest idea (which doesn't work)

- Take the sum of all the residuals.
- Some residuals are negative.

The next simplest idea (which does work)

- Take the square of each residual, and add up those squares.
- This is called the *sum of squares*.


A detour into numerical optimization

A line plot of a quadratic equation

```
x = np.arange(-4, 5, 0.1)
y = x ** 2 - x + 10

xy_data = pd.DataFrame({"x": x,
                        "y": y})

sns.lineplot(x="x",
             y="y",
             data=xy_data)
```



Using calculus to solve the equation

$$y = x^2 - x + 10$$

$$\frac{\partial y}{\partial x} = 2x - 1$$


$$0 = 2x - 1$$

$$x = 0.5$$

$$y = 0.5^2 - 0.5 + 10 = 9.75$$

- Not all equations can be solved like this.
- You can let Python figure it out.

Don't worry if this doesn't make sense, you won't need it for the exercises.



minimize()

```
from scipy.optimize import minimize
```

```
def calc_quadratic(x):  
    y = x ** 2 - x + 10  
    return y
```

```
minimize(fun=calc_quadratic,  
        x0=3)
```

```
    fun: 9.75  
hess_inv: array([[0.5]])  
    jac: array([0.])  
message: 'Optimization terminated successfully.'  
    nfev: 6  
    nit: 2  
    njev: 3  
    status: 0  
success: True  
     x: array([0.49999998])
```

A linear regression algorithm

Define a function to calculate the sum of squares metric.

```
def calc_sum_of_squares(coeffs):  
    intercept, slope = coeffs  
    # More calculation!
```

Call `minimize()` to find coefficients that minimize this function.

```
minimize(  
    fun=calc_sum_of_squares,  
    x0=0  
)
```

Let's practice!

INTERMEDIATE REGRESSION WITH STATSMODELS IN PYTHON